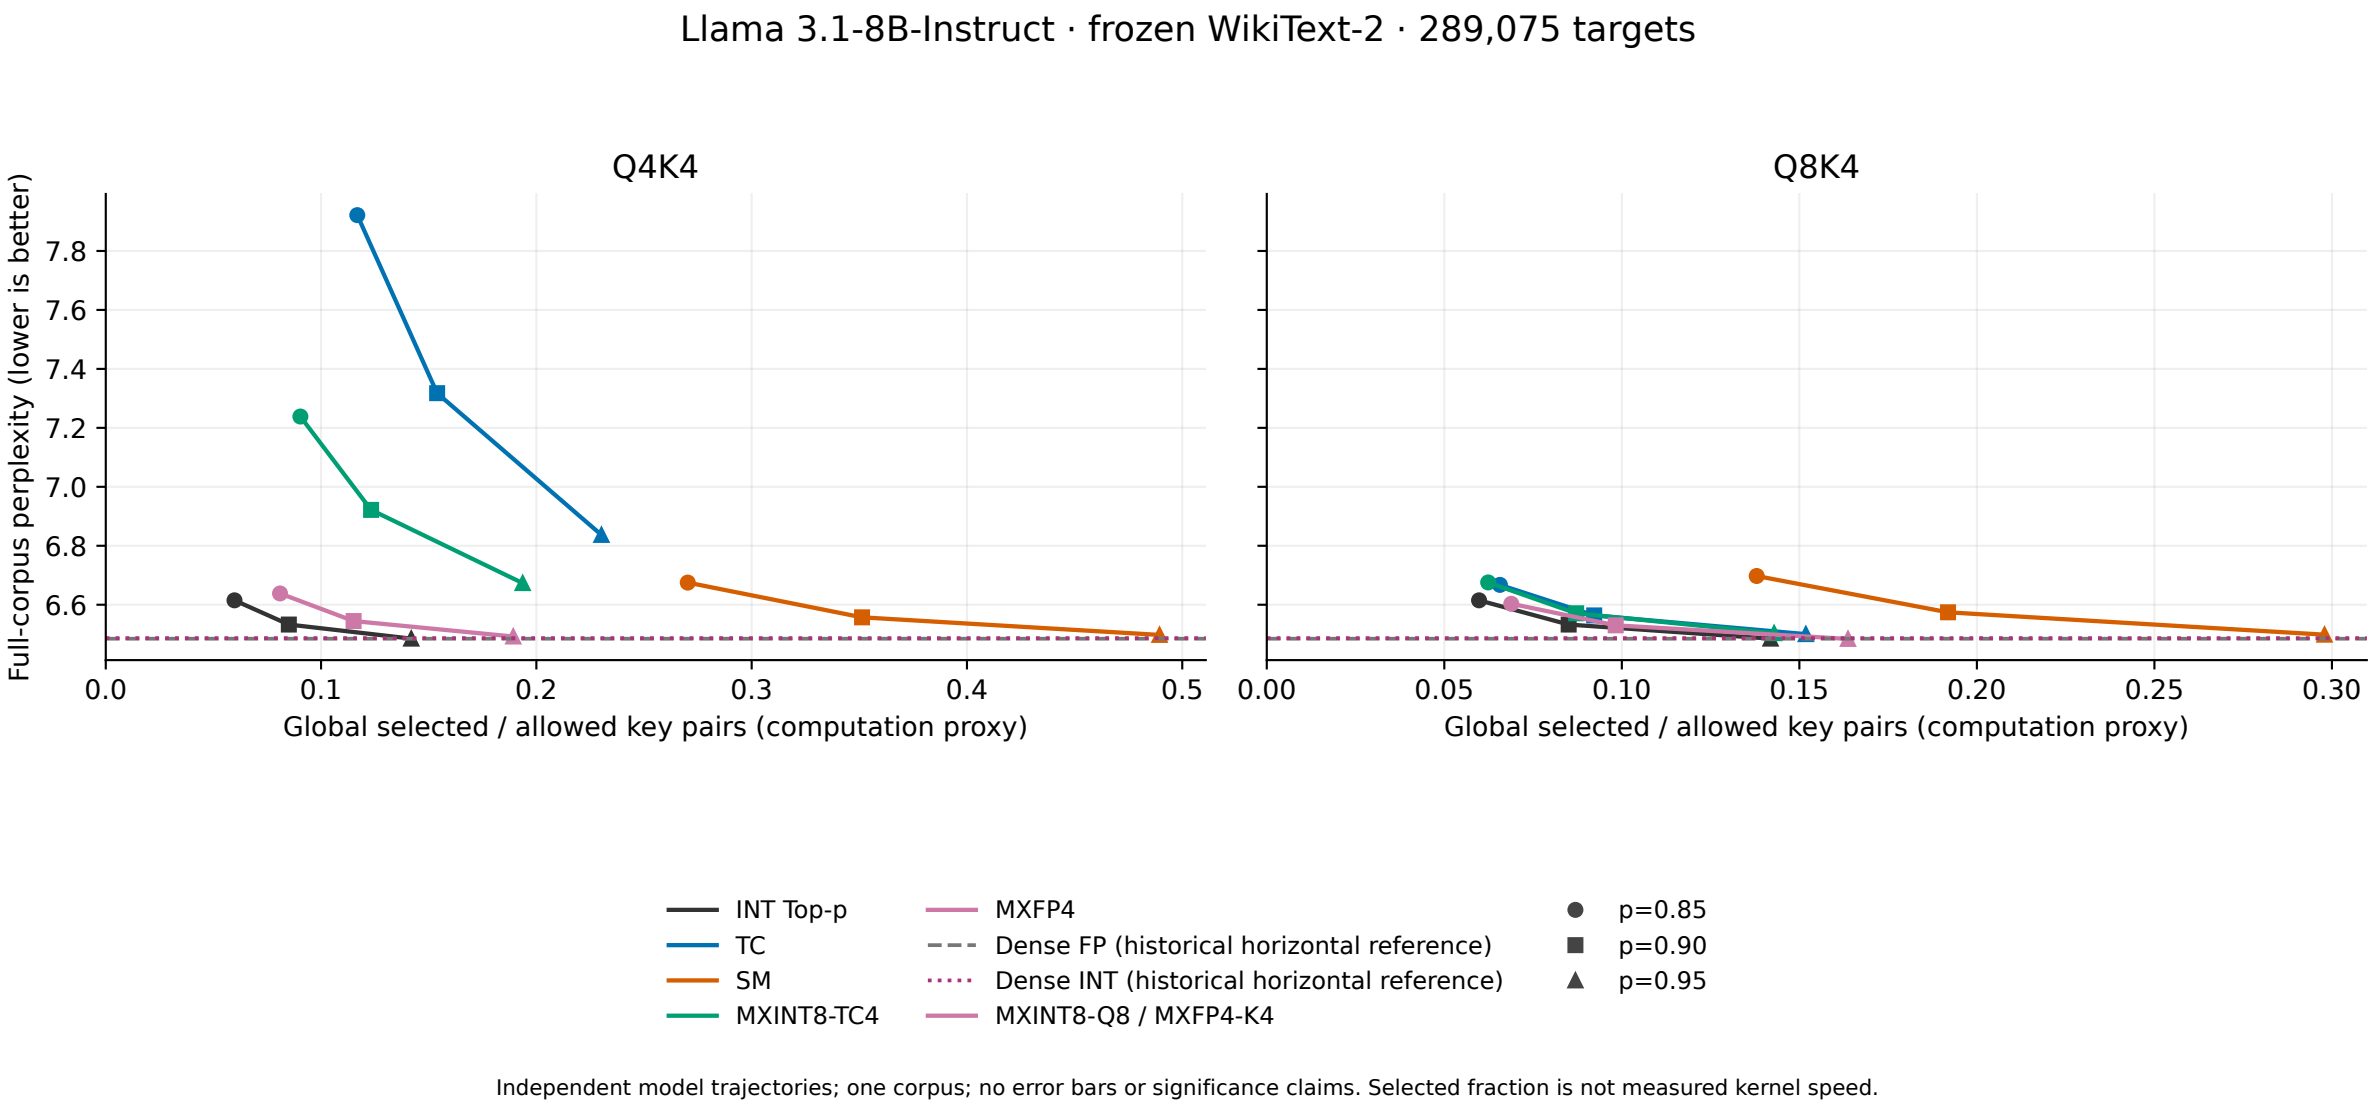
Llama 3.1-8B-Instruct · frozen WikiText-2 · 289,075 targets
Q4K4
Q8K4
Full-corpus perplexity (lower is better)
7.8
7.6
7.4
7.2
7.0
6.8
6.6
0.0
0.1
0.2
0.3
0.4
0.5
0.00
0.05
0.10
0.15
0.20
0.25
0.30
Global selected / allowed key pairs (computation proxy)
Global selected / allowed key pairs (computation proxy)
INT Top-p
TC
SM
MXINT8-TC4
MXFP4
Dense FP (historical horizontal reference)
Dense INT (historical horizontal reference)
MXINT8-Q8 / MXFP4-K4
p=0.85
p=0.90
p=0.95
Independent model trajectories; one corpus; no error bars or significance claims. Selected fraction is not measured kernel speed.